



Total Members	1		
Members	Name	Role	Email
	Harrie	Founder / Solo Builder	nnakastanley01@gmail.com

Overview

Project	Hanko
Track	Track 1 (Autonomous Agents) + Track 3 (New Skills - exit_liquidity)
Problem	Agents that trade on evidence are easy to fake and hard to trust: a verdict with no audit trail, no commitment to what would prove it wrong, and no honest handling of a data source that drops.
Solution	decide() is pure and replayable: every decision hashes its commitment (verdict, size, falsifiers) into its own ID, so replaying stored evidence must reach the same verdict or the mismatch is caught. Missing data shrinks size mechanically. A review loop grades HELD/FALSIFIED against that commitment, not profit.
Key Features	Falsifiers hashed into the decision ID; quality-weighted sizing under missing data; independent-voice detection (echoes don't count twice); review grades against the original commitment, not profit; exit_liquidity (new Track 3 skill); live MCP+REST to RYO's six tools; 198 offline tests.
Target Users	Builders and traders who want an agent's reasoning to be auditable rather than taken on faith - and the judges reviewing this agent, who can replay any decision from the stored bytes it saw.
Scope	A full evidence-to-decision-to-review pipeline: source adapters (X, RSS, RYO's six MCP/REST tools), a pure decision engine, a review/calibration loop, and one new Track 3 skill (exit_liquidity). Live-verified against the real RYO MCP endpoint on 2026-09-05.
Limitations	No live sweep has yet run long enough to produce a calibration curve from real outcomes. exit_liquidity's liquidity_usd fact depends on deep_analysis's optional token_profile lane, unavailable on every live call made - it returns null rather than a guess when that happens. X evidence is demonstrated against fixtures by default; the transport itself is live-tested.

Tech Stack

Frontend Stack	None (static HTML/CSS/vanilla JS, no framework or build step)
Backend Stack	Python 3.11, standard library + httpx for HTTP. No web server - a CLI (hanko) and library, not a hosted service.
AI Model(s)	None at decision time by design: decide() is pure with no model call, which is what makes replay guaranteed. The shipped evidence interpreter is rule-based, not an LLM. Claude (Anthropic) was used as the coding assistant throughout the build.
Other Technologies	httpx (HTTP client), pytest (198 tests), Pillow (site asset generation), Vercel (static hosting + live exit_liquidity endpoint)

Repository / Demo

Github Repository	https://github.com/lamHarrie-Labs/hanko
Demo Video	https://drive.google.com/file/d/1nVfAE-yO4b4ni2CXKxa0DtAsWnY5Zpfx/view
Documentation	README.md (repo root); full technical writeup at tryhanko.vercel.app/docs ; live exit_liquidity demo at tryhanko.vercel.app/try ; exit_liquidity's own standalone README at hanko/skills/exit_liquidity/README.md

Testing Information

How to run the Project	CLI tool (hanko) plus pytest test suite - see commands below
Prerequisites	Python 3.11+; pip
Installation Steps	1) Clone the repo 2) pip install -e ".[dev]" 3) Copy .env.example to .env and fill in credentials
Environment Variables	XAI_API_KEY (X evidence); RYO_MCP_URL, RYO_MCP_KEY (RYO platform, MCP transport); RYO_API_BASE, RYO_API_KEY (optional REST fallback)
Build Command	None for the package. Site: python site/build.py
Run Command	hanko decide x --token SOL --subject <handle>; hanko sweep --watchlist watchlist.json; hanko exit-liquidity SOL --size 50000; hanko audit

Test Command	pytest
Test Account(s)	None required - all RYO tools are read-only research APIs authenticated by an environment-variable credential, not a login.